

**Project Name:** Car Price Prediction

**Author:** Laxmikant walzade

## • Problem Definition.

- With the covid 19 impact in the market, we have seen lot of changes in the car market. Now some cars are in demand hence making them costly and some are not in demand hence cheaper. One of our clients works with small traders, who sell used cars. With the change in market due to covid 19 impact, our client is facing problems with their previous car price valuation machine learning models. So, they are looking for new machine learning models from new data. We have to make car price valuation model.
- This dataset include following columns:-
- 'Price', 'Name', 'Kilometers\_Driven', 'Last\_Service', 'Registration', 'Registered\_in', 'Fuel\_Type', 'Transmission', 'Insurance', 'Airbags', 'Seat\_Upholstery', 'Integrated\_Music', 'Rear\_View\_Mirrors', 'Engine\_start\_stop', 'Central\_Locking', 'Sunroof\_Moonroof', 'Rear\_AC', 'Power\_Windows', 'Headlamps', 'Fuel\_type', 'Engine\_type', 'Drivetrain', 'Mileage', 'Steering\_type', 'Transmission\_type', 'Max\_power', 'Fuel\_tank\_capacity', 'Seating\_capacity', 'Alternate\_fuel\_type', 'History'

## • Data Analysis.

### DATASET

	Price	Name	Kilometers_Driven	Last_Service	Registration	Registered_in	Fuel_Type	Transmission	Insurance	Airbags	...
0	Fixed Price ₹3,37,599	2019 Maruti Alto 800 LXI MANUAL	79,205 km	79,205km (28 Aug 2022)	GJ-17-x-xxxx	May-19	Petrol + CNG	MANUAL	Valid upto May 55660 3rd Party	NaN	...
1	Fixed Price ₹10,56,899	2020 Maruti S Cross ZETA AT 1.5 SHVS	11,707 km	11,707km (27 Jul 2022)	WB-06-x-xxxx	Dec-20	Petrol	NaN	Valid upto May 55660 3rd Party	-	...
2	Fixed Price ₹15,89,699	2019 Honda Civic VX CVT i-VTEC	13,878 km	13,878km (18 Jul 2022)	MH-02-x-xxxx	NaN	Petrol	NaN	Valid upto May 55660 3rd Party	4 Airbags (Driver, Front Passenger, Driver Side)	...
3	Fixed Price ₹18,79,099	2020 MG HECTOR PLUS SHARP DCT	11,086 km	11,086km (29 Aug 2022)	MH-14-x-xxxx	Aug-20	Petrol	NaN	Valid upto May 55660 3rd Party	6 Airbags (Driver, Front Passenger, 2nd Row, 2nd Row, 2nd Row, 2nd Row)	...
4	Fixed Price ₹7,08,299	2018 Maruti Swift ZXI AMT AUTOMATIC	41,249 km	41,249km (08 Sep 2022)	TN-14-x-xxxx	Aug-18	Petrol	AUTOMATIC	Valid upto May 55660 3rd Party	2 Airbags (Driver, Front Passenger)	...

- We have total 30 columns including the label i.e Price column.

- **Pre-Processing Steps**

1. Identifying sources of the data
2. Analysing the information
3. Cleaning and handling the information
4. Selecting the most significant elements
5. Writing down findings and observations
6. Using various models to train the data
7. Selecting the best-fitted model for predictions
8. Predicting results for test information

- **Pre-Processing Pipeline.**

- First let's check the data type of the dataset.

```
: ▶ df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4409 entries, 0 to 4408
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Price                                4409 non-null   object
1   Name                                4409 non-null   object
2   Kilometers_Driven                   4409 non-null   object
3   Last_Service                        4409 non-null   object
4   Registration                         4409 non-null   object
5   Registered_in                       3844 non-null   object
6   Fuel_Type                           4409 non-null   object
7   Transmission                        4279 non-null   object
8   Insurance                           4409 non-null   object
9   Airbags                             2290 non-null   object
10  Seat_Upholstery                      900 non-null    object
11  Integrated_Music                     3580 non-null   object
12  Rear_View_Mirrors                    362 non-null    object
13  Engine_start_stop                    1502 non-null   object
14  Central_Locking                      3925 non-null   object
15  Sunroof_Moonroof                     681 non-null    object
16  Rear_AC                              1328 non-null   object
17  Power_Windows                        4257 non-null   object
18  Headlamps                            4125 non-null   object
19  Fuel_type                            4409 non-null   object
20  Engine_type                          4190 non-null   object
21  Drivetrain                           4300 non-null   object
22  Mileage                              4349 non-null   object
23  Steering_type                        4170 non-null   object
24  Transmission_type                    4409 non-null   object
25  Max_power                            4409 non-null   object
26  Fuel_tank_capacity                   4397 non-null   object
27  Seating_capacity                     4409 non-null   object
28  Alternate_fuel_type                  368 non-null    object
29  History                              4409 non-null   object
dtypes: object(30)
memory usage: 1.0+ MB
```

All columns are Object type data which needs to change to Integer, since it is important for model building as model does not consider the string value.

- There are no Null Values present in the dataset so we can move further.

```
▶ for i in df.columns:  
    a = df[i].isna().sum()  
    if a > 0:  
        print(i, 'column has', a, 'NaN values')
```

```
Registered_in column has 565 NaN values  
Transmission column has 130 NaN values  
Airbags column has 2119 NaN values  
Seat_Upholstery column has 3509 NaN values  
Integrated_Music column has 829 NaN values  
Rear_View_Mirrors column has 4047 NaN values  
Engine_start_stop column has 2907 NaN values  
Central_Locking column has 484 NaN values  
Sunroof_Moonroof column has 3728 NaN values  
Rear_AC column has 3081 NaN values  
Power_Windows column has 152 NaN values  
Headlamps column has 284 NaN values  
Engine_type column has 219 NaN values  
Drivetrain column has 109 NaN values  
Mileage column has 60 NaN values  
Steering_type column has 239 NaN values  
Fuel_tank_capacity column has 12 NaN values  
Alternate_fuel_type column has 4041 NaN values
```

- Sometimes some unwanted things can be found in a dataset which are equivalent to Null values. It is important to take care of such cases.

- Need to clean and drop some columns

- dropped 'Registered\_in', 'Airbags', 'Seat\_Upholstery', 'Integrated\_Music', 'Rear\_View\_Mirrors', 'Engine\_start\_stop', 'Rear\_AC' and 'Alternate\_fuel\_type' column as it has a lot of NaNs as it may affect the result of the model

```
] df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4409 entries, 0 to 4408
Data columns (total 22 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Price                 4409 non-null   object
 1   Name                  4409 non-null   object
 2   Kilometers_Driven     4409 non-null   object
 3   Last_Service          4409 non-null   object
 4   Registration          4409 non-null   object
 5   Fuel_Type             4409 non-null   object
 6   Transmission          4279 non-null   object
 7   Insurance             4409 non-null   object
 8   Central_Locking       3569 non-null   object
 9   Sunroof_Moonroof      4053 non-null   object
10  Power_Windows         3900 non-null   object
11  Headlamps             3768 non-null   object
12  Fuel_type              4409 non-null   object
13  Engine_type           3836 non-null   object
14  Drivetrain            3946 non-null   object
15  Mileage               3987 non-null   object
16  Steering_type         3816 non-null   object
17  Transmission_type     4055 non-null   object
18  Max_power             4055 non-null   object
19  Fuel_tank_capacity     4043 non-null   object
20  Seating_capacity      4055 non-null   object
21  History               4408 non-null   object
dtypes: object(22)
memory usage: 757.9+ KB
```

- Filling the nan values

```
df['Transmission'] = df['Transmission'].fillna('MANUAL')
df['Sunroof_Moonroof'] = df['Sunroof_Moonroof'].fillna('No')
df['Power_Windows'] = df['Power_Windows'].fillna('Front & Rear')
df['Drivetrain'] = df['Drivetrain'].fillna('FWD')
df['Mileage'] = df['Mileage'].fillna(df['Mileage'].mean())
df['Transmission_type'] = df['Transmission_type'].fillna('Manual')
df['Max_power'] = df['Max_power'].fillna(df['Max_power'].mean())
df['Fuel_tank_capacity'] = df['Fuel_tank_capacity'].fillna(df['Fuel_tank_capacity'].mean())
df['Seating_capacity'] = df['Seating_capacity'].fillna(df['Seating_capacity'].mean())
df['Steering_type'] = df['Steering_type'].fillna('Power assisted (Electric)')
```

Filled NaN with Mode in case of categorical data and Mean in case of continuous data

We have only 18 columns left after EDA

- Data looks fine to encode into Int

```
In [83]: df[['Price1', 'Price']] = df['Price'].str.split('\n', expand=True)

In [85]: df = df.drop('Price1', axis=1)

In [86]: df1 = df['Name'].str.split(' ', expand=True)

In [89]: df['Launch_Year'] = df1[0]

In [91]: df[['Kilometers_Driven', 'test']] = df['Kilometers_Driven'].str.split('km', expand=True)

In [93]: df = df.drop('test', axis=1)

In [94]: df[['Registration', 'Registration_code', 'test', 'test']] = df['Registration'].str.split('-', expand=True)

In [96]: df['Registration'] = df['Registration'] + '-' + df['Registration_code']

In [98]: df = df.drop(['Registration_code', 'test'], axis=1)

In [100]: df[['Last_Service', 'test']] = df['Last_Service'].str.split('km', expand=True)

In [102]: df = df.drop('test', axis=1)

In [104]: df.to_csv('CarData1_Final.csv')

In [109]: df = pd.read_csv('CarData1_Final.csv')
df.head()

for i in df.columns:
    if df[i].dtypes == 'object':
        df[i] = enc.fit_transform(df[i].values.reshape(-1, 1))
```

Encoded data into Int

## • Feature Engg

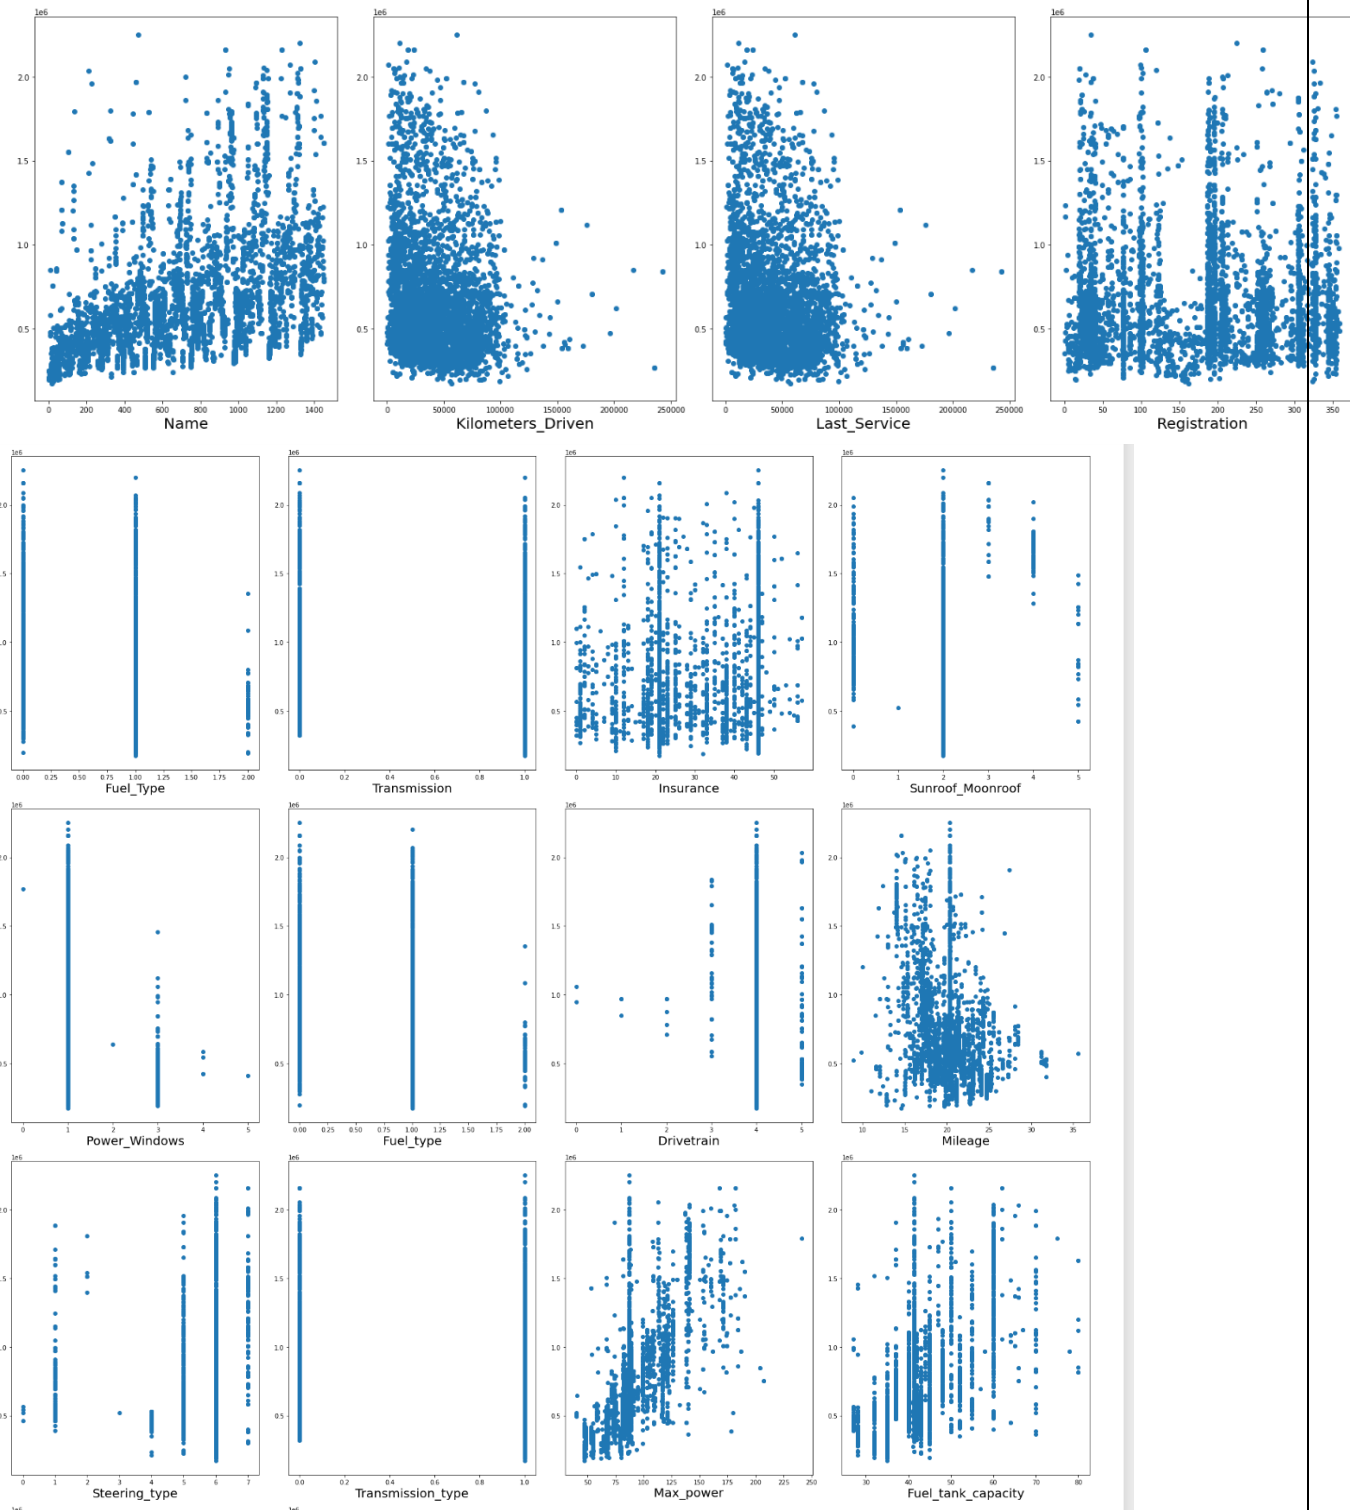
```
x = df.drop(['Price'],axis=1)
y = df['Price']
```

```
plt.figure(figsize=(25,40), facecolor='white')
```

```
plotno = 1
```

```
for column in x:
    if plotno <= 18:
        ax = plt.subplot(5,4,plotno)
        plt.scatter(x[column],y)
        plt.xlabel(column,fontsize=20)
```

```
    plotno+=1
plt.tight_layout()
```





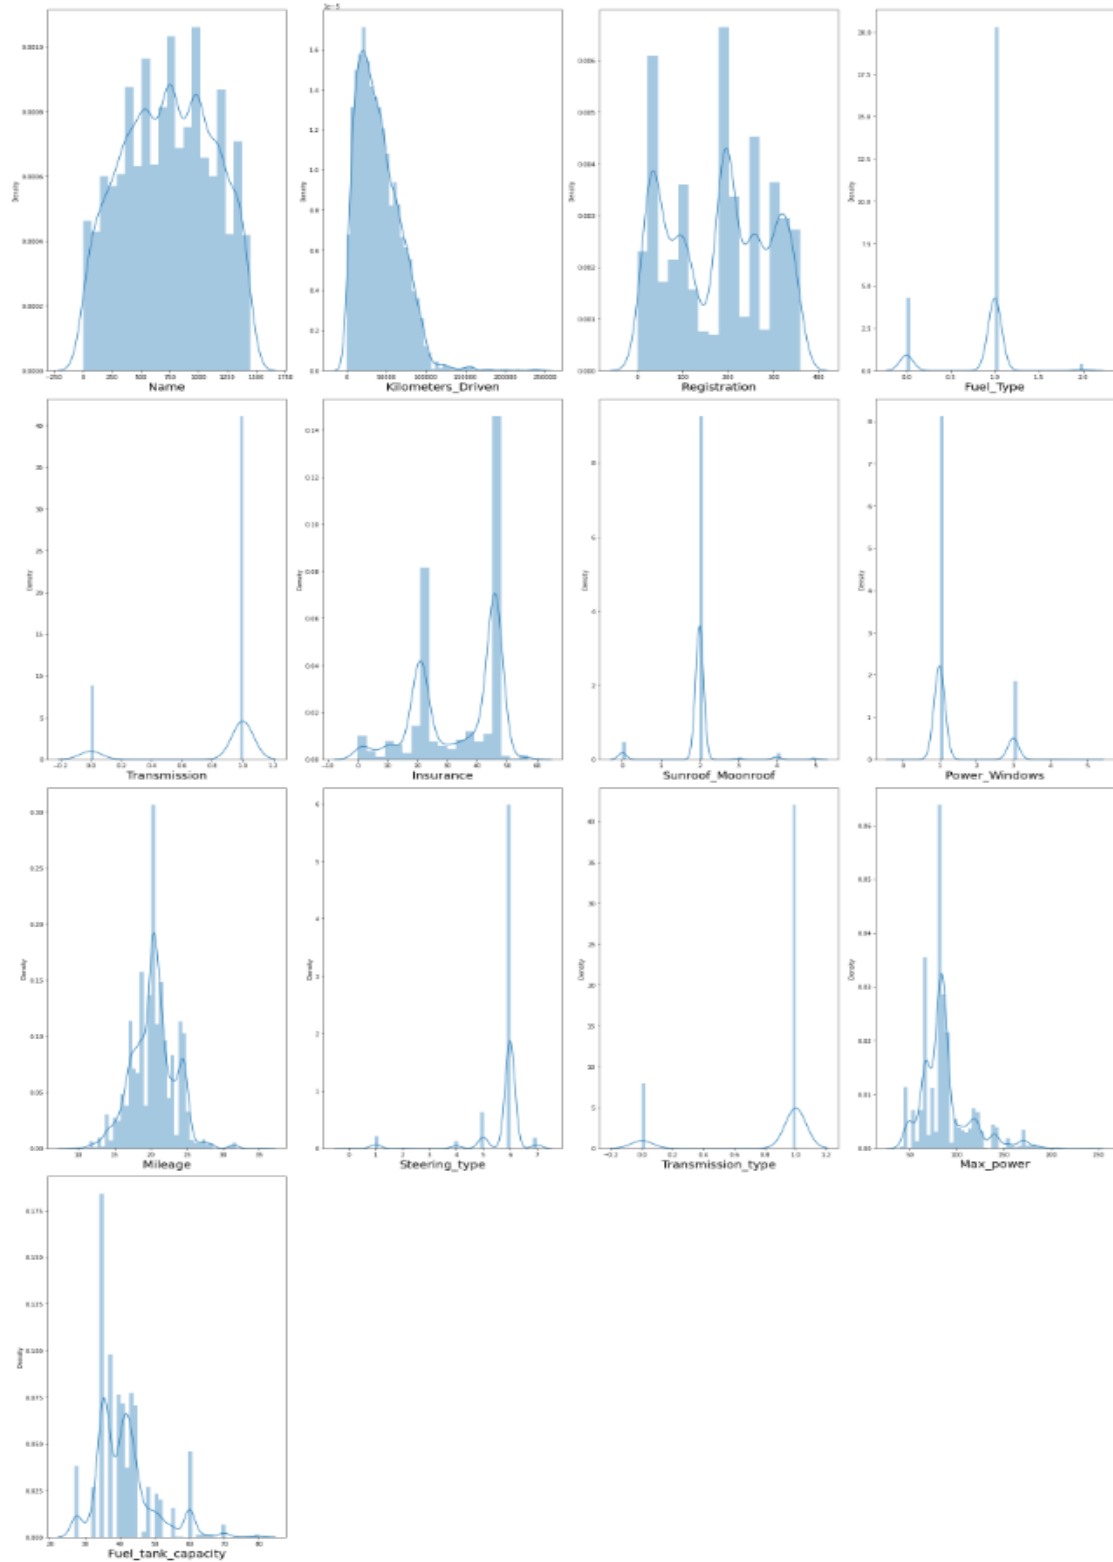
- Checking for skew data

```
plt.figure(figsize=(25,40), facecolor='white')

plotno = 1

for column in x:
    if plotno <= 16:
        ax = plt.subplot(4,4,plotno)
        sns.distplot(x[column])
        plt.xlabel(column,fontsize=20)

        plotno+=1
plt.tight_layout()
```



```

▶ for i in df.columns:
    a = df[i].skew()
    print(i, '=', a)

```

```

Price = 1.5283538565208503
Name = -0.04499556646526514
Kilometers_Driven = 1.2745889548113183
Registration = -0.0183436719157604
Fuel_Type = -1.1372243790287218
Transmission = -1.6898684909473625
Insurance = -0.5436849312158353
Sunroof_Moonroof = -0.3684005541026811
Power_Windows = 1.6248449443874342
Mileage = 0.08432065037061806
Steering_type = -3.9960390541366597
Transmission_type = -1.8634271510542266
Max_power = 1.454121727378636
Fuel_tank_capacity = 1.1907958191669374

```

Not considering skewness of categorical data columns


- Kilometers\_Driven, Max\_power and Fuel\_tank\_capacity has skewness

Database is ready to remove outliers if exist

```

▶ a = x.columns.values
col = 35
row = 30
plt.figure(figsize = (col,3*row))
for i in range(0, len(a)):
    plt.subplot(row,col,i+1)
    sns.boxplot(data = x[a[i]],color='blue',orient='v')
plt.tight_layout()

```



LotFrontage, LotArea, BsmtFinSF1, BsmtUnfSF, TotalBsmtSF, 1stFlrSF, LowQualFinSF, GrLivArea, GarageArea, WoodDeckSF, OpenPorchSF, EnclosedPorch, ScreenPorch and PoolArea Column has outliers

- Skewness removed by quantile method

- **Scaling the data**

- ▶ 

```
x = df.drop(['Price'],axis=1)  
y = df['Price']
```

- ▶ 

```
scaler = StandardScaler()  
X_scale = scaler.fit_transform(x)
```

-

- **EDA Concluding Remark.**

1. The information was not organized and coordinated and subsequently cleaned the information utilizing different information cleaning and pre-handling techniques.
2. There are numerous anomalies present in the information consequently eliminating exceptions
3. There was a skewness in the information thus have eliminated the skewness from the information.
4. There was an irregularity in the information thus have utilized SMOTE strategy to balance the information.
5. Scaled the data utilizing Standard Scalar to make the information normalized to fabricate a model.

## • Hardware and Software Requirements and Tools Used

### 1. Libraries and packages used

- import numpy as np - For Numpy work
- import pandas as pd - To work on DataFrame
- import seaborn as sns - Plotting Graphs
- import matplotlib.pyplot as plt - Plotting Graphs
- import pickle – To save the Model
- from sklearn.preprocessing import StandardScaler (To scale the train data), OrdinalEncoder(To encode object data to Integer), PowerTransformer (To remove skewness from dataset)
- from statsmodels.stats.outliers\_influence import variance\_inflation\_factor
- enc = OrdinalEncoder() = Assigned OrdinalEncoder to variable
- from sklearn.model\_selection import train\_test\_split, GridSearchCV, cross\_val\_score,(To split the data into train and test, Search the best parameters, to calculate cross validation score)
- from sklearn.metrics import accuracy\_score, classification\_report, confusion\_matrix, roc\_curve, roc\_auc\_score - To calculate and analyse model metrics.
- from sklearn import metrics

### Models which are used

- from sklearn.ensemble import RandomForestClassifier
- from sklearn.linear\_model import LogisticRegression
- from sklearn.ensemble import GradientBoostingClassifier
- from sklearn.tree import DecisionTreeClassifier
- from sklearn.neighbors import KNeighborsClassifier
- from sklearn.svm import SVC
  
- import warnings
- warnings.filterwarnings('ignore') - To ignore unwanted Warnings

### 2. Hardware used – 11th Gen Intel(R) Core (TM) i3-1115G4 @ 3.00GHz 3.00 GHz with 8.00 GB RAM and Windows 11

### 3. Software used – Anaconda and Jupyter Notebook to build the model.

## • Building Machine Learning Models.

I have built 6 machine learning models to predict the label. Below are the machine learning models which are been used.

1. LogisticRegression
2. RandomForestClassifier
3. DecisionTreeClassifier
4. GradientBoostingClassifier
5. Support Vector Classifier
6. KNeighborsClassifier

### 1. LogisticRegression:

Have used “For Loop” to find out the highest accuracy score with different random state ranging from 0-100 and using that random state to split the data into train and test data.

```
] | ▶ reg = LinearRegression()
    for i in range(0,100):
        x_train,x_test,y_train,y_test = train_test_split(X_scale,y,test_size = 0.25,random_state = i)
        reg.fit(x_train, y_train)
        x_pred = reg.predict(x_train)
        y_pred = reg.predict(x_test)
        print("At Random state", (i), "the training accuracy is :-", (r2_score (y_train,x_pred)))
        print("At Random state", (i), "the testing accuracy is :-", (r2_score (y_test,y_pred)))
        print('\n')
At Random state 24 the testing accuracy is :- 0.6722909420571519

At Random state 25 the training accuracy is :- 0.6726468395656839
At Random state 25 the testing accuracy is :- 0.6746602758572529

At Random state 26 the training accuracy is :- 0.6634338773478726
At Random state 26 the testing accuracy is :- 0.704169495101693

At Random state 27 the training accuracy is :- 0.6695078428641499
At Random state 27 the testing accuracy is :- 0.6858194397585636

At Random state 28 the training accuracy is :- 0.6681201936149961
At Random state 28 the testing accuracy is :- 0.6908866202508435

At Random state 29 the training accuracy is :- 0.6671114526023241
```

Have used “Define Function” to define a machine learning model code that automatically provides the train and test accuracy code.

Formulas:

$y\_pred = clf.predict(x\_train)$  = Predicting train data

$accuracy\_score(y\_train, y\_pred)$  = Calculating train accuracy score (comparing  $y\_pred$  data with  $y\_train$  data)

pred = clf.predict(x\_test) = Predicting test data

accuracy\_score(y\_test, pred) = Calculating test accuracy score (comparing pred data with y\_test data)

```
] : ➤ x_train,x_test,y_train,y_test = train_test_split(X_scale,y,test_size = 0.25,random_state = 66)
```

```
] : ➤ def print_score(clf, x_train,x_test,y_train,y_test, train=True):  
    if train:  
        y_pred = clf.predict(x_train)  
  
        print('\n=====Train Result=====')  
        print(f'Accuracy Score: {r2_score (y_train,y_pred)*100:.2f}%')  
  
    elif train==False:  
        pred = clf.predict(x_test)  
  
        print('\n=====Test Result=====')  
        print(f'Accuracy Score: {r2_score (y_test,pred)*100:.2f}%')  
  
        print ('\n mean_absolute_error',mean_absolute_error(y_test,pred))  
        print ('\n mean_squared_error',mean_squared_error (y_test,pred))
```

Trained the data and run the “Def” function

```
] : ➤ reg = LinearRegression()  
    reg.fit(x_train,y_train)  
  
    print_score(reg,x_train,x_test,y_train,y_test, train=True)  
    print_score(reg,x_train,x_test,y_train,y_test, train=False)
```


```
=====Train Result=====  
Accuracy Score: 67.35%
```

```
=====Test Result=====  
Accuracy Score: 67.33%
```

```
mean_absolute_error 102376.51939422284
```

```
mean_squared_error 28665347329.86425
```

```
] : ➤ plt.scatter(y_test, y_pred)  
    plt.xlabel('Actual Quality')  
    plt.ylabel('Predicted Quality')  
    plt.title('Actual vs Predicted')  
    plt.show()
```



## 2. RandomForestClassifier:

```
[62]: ▶ rfr = RandomForestRegressor()  
      rfr.fit(x_train,y_train)  
  
      print_score(rfr,x_train,x_test,y_train,y_test, train=True)  
      print_score(rfr,x_train,x_test,y_train,y_test, train=False)
```

```
=====Train Result=====  
Accuracy Score: 97.84%
```

```
=====Test Result=====  
Accuracy Score: 86.79%
```

```
mean_absolute_error 54218.23786166842
```

```
mean_squared_error 11587753426.816917
```

## 3. DecisionTreeClassifier:

```
]▶ dtr = DecisionTreeRegressor()  
   dtr.fit(x_train,y_train)  
  
   print_score(dtr,x_train,x_test,y_train,y_test, train=True)  
   print_score(dtr,x_train,x_test,y_train,y_test, train=False)
```

```
=====Train Result=====  
Accuracy Score: 100.00%
```

```
=====Test Result=====  
Accuracy Score: 76.61%
```

```
mean_absolute_error 61022.305174234425
```

```
mean_squared_error 20520848167.263992
```



#### 4. GradientBoostingClassifier:

```
: ▶ gbdtr = GradientBoostingRegressor()
    gbdtr.fit(x_train,y_train)

print_score(gbdtr,x_train,x_test,y_train,y_test, train=True)
print_score(gbdtr,x_train,x_test,y_train,y_test, train=False)

=====Train Result=====
Accuracy Score: 85.11%

=====Test Result=====
Accuracy Score: 79.90%

mean_absolute_error 75192.5135144641

mean_squared_error 17635591320.572407
```

#### 5. Support Vector Classifier:

```
: ▶ svr = SVR()
    svr.fit(x_train,y_train)

print_score(svr,x_train,x_test,y_train,y_test, train=True)
print_score(svr,x_train,x_test,y_train,y_test, train=False)

=====Train Result=====
Accuracy Score: -5.74%

=====Test Result=====
Accuracy Score: -4.99%

mean_absolute_error 205857.07725353728

mean_squared_error 92104971953.68492
```

#### 6. KNeighborsClassifier:

```
▶ knr = KNeighborsRegressor()
  knr.fit(x_train,y_train)

print_score(knr,x_train,x_test,y_train,y_test, train=True)
print_score(knr,x_train,x_test,y_train,y_test, train=False)

=====Train Result=====
Accuracy Score: 82.39%

=====Test Result=====
Accuracy Score: 71.56%

mean_absolute_error 89830.3626187962

mean_squared_error 24947042897.9007
```

- **Findings**

- LinearRegression train accuracy score 67.35% and test accuracy score 67.33%
- Support Vector Regression train accuracy score -5.74% and test accuracy score -4.99%
- DecisionTreeRegressor train accuracy score 100.00% and test accuracy score 76.61%
- AdaBoostRegressor train accuracy score 57.60% and test accuracy score 55.07%
- GradientBoostingRegressor train accuracy score 85.11% and test accuracy score 79.90%
- RandomForestRegressor train accuracy score 97.84% and test accuracy score 86.79%
- KNeighborsRegressor train accuracy score 82.39% and test accuracy score 71.56%

- **Model Selection:**

LinearRegression it has low variance between train and test result and has high 67.35% and 67.33% accuracy i.e., respectively.